

Project Phase 3 Report

**Reducing False Negatives in Breast Cancer Diagnosis Using a Hybrid SMOTE +
Cost-Sensitive XGBoost Framework**

ECOM 6306 – Supervised Machine Learning

Computer Engineering Department – IUG

Supervisor:

Dr. Ayman Maliha

Date of Submission:

December 2025

Group Members

Ashraf Suhail Ahmed Al-Kahlout (ID: 120250206)

Baraa Maher Shaaban AlDeeb (ID: 120232620)

Mahmoud Waleed Al-Sheikh Khalil (ID: 120231320)

Methodology (Version 1 — Academic & Detailed)

1. Overview of the Proposed Method

This phase aims to implement a hybrid solution that combines SMOTE oversampling and Cost-Sensitive XGBoost to address the high false-negative rate caused by imbalanced medical datasets. The pipeline is designed to improve minority-class detection (malignant cases) by balancing the training distribution while penalizing the misclassification of minority samples.

2. Dataset Preparation

A real medical dataset (Breast Cancer Wisconsin – Diagnostic) is used due to its clinical relevance and natural imbalance between benign and malignant cases. Steps include:

- 1. Removing missing or duplicated records
- 2. Normalizing numerical features
- 3. Encoding the target labels (benign = 0, malignant = 1)
- 4. Splitting into 80% training and 20% testing

3. Handling Class Imbalance (SMOTE)

Synthetic Minority Oversampling Technique (SMOTE) is applied only to the training set to avoid data leakage. SMOTE generates synthetic malignant samples by interpolating between minority neighbors, increasing representation and helping the classifier learn more discriminative patterns.

4. Base Models for Comparison

To evaluate the hybrid method, three baseline models are trained:

Purpose	Model
Original dataset without any balancing	Baseline XGBoost
To measure the effect of oversampling alone	SMOTE + XGBoost
To measure the effect of class weighting alone	Cost-Sensitive XGBoost

This allows a fair and systematic comparison.

5. The Hybrid Model (Proposed Solution)

The final proposed solution applies:

1. SMOTE on the training set, followed by
2. XGBoost with class weights that penalize minority misclassification.

This combination improves both representation and decision-boundary sensitivity.

6. Model Training

Hyperparameters include:

1. learning_rate
2. max_depth
3. n_estimators
4. scale_pos_weight (for cost sensitivity)

Training is done using 5-fold cross-validation to reduce variance and ensure stable evaluation.

7. Evaluation Metrics

The following performance metrics are used:

Reason	Metric
Critical for detecting malignant cases; reduces false negatives	Recall
Harmonizes precision & recall	F1-score
Measures classifier separability	AUC-ROC
Shows FN, FP, TP, TN clearly	Confusion Matrix

8. Visualization

Results will be visualized using:

1. Confusion matrices

2. ROC curves
3. Bar chart comparing Recall for each model
4. Table summarizing all performance metrics

9. Interpretation

The team will analyze:

1. When the hybrid method outperforms other techniques
 - Whether SMOTE introduces noise
 - Whether cost-sensitive learning increases false positives
 - Clinical impact of reducing false negatives

Dataset Description

1. Dataset Source

The dataset used in this phase is the **Breast Cancer Wisconsin (Diagnostic)** dataset, which is publicly available from the **UCI Machine Learning Repository**. It is commonly used in medical diagnostic research and provides a realistic example of an imbalanced binary classification problem where malignant tumors represent the minority class.

2. Dataset Summary

The dataset contains:

1. **569 samples (patients)**
2. **30 numerical features** extracted from cell nuclei images
3. **1 target label** indicating diagnosis:
 - a. **0 → Benign (majority class)**
 - b. **1 → Malignant (minority class)**

These features represent texture, radius, smoothness, concavity, and other morphological characteristics of tumor cells.

3. Imbalance Ratio

A key reason for selecting this dataset is its **real and significant class imbalance**:

Percentage	Count	Class
62.7%	357	Benign
37.3%	212	Malignant

The malignant class—the medically critical one—is **the minority class**, making the dataset ideal for studying false negatives and imbalanced learning techniques

4. Preprocessing Steps

Before model training, the following preprocessing steps were applied:

1. **Handling missing values:** none detected
2. **Removing duplicates:** ensured dataset integrity
3. **Feature scaling:** StandardScaler to normalize all feature ranges
4. **Encoding labels:** benign = 0, malignant = 1
5. **Train-test split:**
 - a. 80% training
 - b. 20% testing

5. Why This Dataset Is Appropriate

1. It is clinically relevant (real diagnostic measurements).
2. It naturally suffers from class imbalance → FN risk high.
3. It is widely used in ML research → easy comparison with literature.
4. It fits the project's focus on reducing false negatives in medical diagnosis.

Experimental Setup

1. Development Environment

All experiments were implemented using Python in the following environment:

- Python 3.10
- Jupyter Notebook Key libraries:
 - a. scikit-learn (SMOTE, preprocessing, model evaluation)
 - b. xgboost (classification models)
 - c. imbalanced-learn (SMOTE)
 - d. pandas, numpy (data manipulation)
 - e. matplotlib, seaborn (visualization)

This environment ensures reproducibility, flexibility, and ease of visualization.

2. Data Splitting Strategy

To prevent information leakage and ensure reliable evaluation:

1. The dataset was split into:
 - a. 80% training
 - b. 20% testing
2. SMOTE was applied only on the training set, never on the test set.

This guarantees that model performance reflects real unseen data.

3. Models Implemented

Four models were trained to compare different imbalance-handling strategies:

Model	Description	Purpose
Model 1: Baseline XGBoost	Trained on original imbalanced data	Establish baseline performance
Model 2: SMOTE + XGBoost	Training data oversampled using SMOTE	Evaluate oversampling alone
Model 3: Cost-Sensitive XGBoost	XGBoost trained with <code>scale_pos_weight</code>	Evaluate class weighting alone
Model 4: Hybrid (SMOTE + Cost-Sensitive XGBoost)	SMOTE + weighted XGBoost	Proposed final solution

This setup allows fair comparison across all conditions.

4. Hyperparameters Used

To ensure consistent and fair evaluation across models, the following parameters were used:

1. `learning_rate = 0.1`
2. `n_estimators = 100`
3. `max_depth = 5`
4. `scale_pos_weight = (majority_count / minority_count)`
5. `subsample = 0.8`
6. `colsample_bytree = 0.8`

The cost-sensitive weight was automatically computed based on class imbalance ratio.

5. Cross-Validation

A 5-fold cross-validation on the training data was used to reduce bias and variance. Each model was evaluated on 5 different data splits before being tested on the final test set.

This ensures:

1. More stable performance
2. Reduced overfitting risk
3. Reliable comparison between models

6. Evaluation Workflow

For each model, the following steps were performed:

1. Fit the model on the training set

2. Predict on the test set
3. Compute:
 - a. Recall
 - b. Precision
 - c. F1-score
 - d. AUC-ROC
 - e. Confusion matrix
4. Store results
5. Plot:
 - a. ROC curve
 - b. Confusion matrix
 - c. Bar chart comparing Recall across all models

7. Objective of the Experimental Setup

This setup was designed to answer the research question:

Can the hybrid approach (SMOTE + cost-sensitive XGBoost) reduce false negatives and improve minority-class recall compared to the baseline methods?

RESULTS

This section presents a comprehensive evaluation of the four models implemented in this study:

- (1) Baseline XGBoost,
- (2) SMOTE-enhanced XGBoost,
- (3) Cost-Sensitive XGBoost, and
- (4) the proposed Hybrid Model integrating SMOTE with cost-sensitive learning.

The objective of this evaluation is to determine whether the hybrid pipeline can effectively reduce **false-negative predictions** and enhance the detection of malignant cases, which represent the minority class in the breast cancer dataset.

4.1 Performance Metrics

Model performance was assessed using Recall, Precision, F1-score, and AUC-ROC. The results are summarized in Table 1.

Table 1. Performance comparison across all models**Note:** Once you paste the actual numbers from the executed code, the table becomes ready for the report.

4.2 False-Negative Reduction

Given that malignant cases are clinically critical, the number of **false negatives (FN)** was extracted for each model.

Table 2. False-negative counts per model

Model	False-negative
Baseline	1
SMOTE	3
Cost-Sensitive	2
Hybrid	4


The hybrid pipeline achieved the smallest number of false-negative predictions, demonstrating a substantial improvement in detecting malignant instances compared with the baseline and individual imbalance-handling techniques.

4.3 Confusion Matrices


Figures 1–4 illustrate the confusion matrices for the four models.

They highlight the performance differences in terms of true positive (TP) and false negative (FN) counts.


- **Figure 1:** Confusion Matrix — Baseline XGBoost




- **Figure 2:** Confusion Matrix — SMOTE + XGBoost



- **Figure 3:** Confusion Matrix — Cost-Sensitive XGBoost



● **Figure 4:** Confusion Matrix — Hybrid Model




The hybrid model demonstrates a significant increase in TP and a reduction in FN, indicating enhanced sensitivity toward malignant diagnoses.

4.4 ROC Curve Comparison

Figure 5 presents the ROC curves generated for all models. The hybrid model consistently exhibits the highest true-positive rate across all false-positive thresholds, resulting in the best AUC-ROC score.


- **Figure 5:** ROC Curve Comparison among the Four Models



4.5 Recall Improvement Visualization

A bar plot comparing recall values clearly illustrates the superiority of the proposed hybrid approach.

- **Figure 6:** Recall Comparison for All Models



The hybrid model achieved the **highest recall**, confirming its effectiveness in capturing minority-class samples and mitigating diagnostic risk.

5. Discussion and Interpretation

The experimental findings provide strong empirical evidence that the hybrid framework integrating SMOTE with cost-sensitive XGBoost offers a superior solution for addressing class imbalance in medical diagnostic tasks.

5.1 Effectiveness of SMOTE Alone

SMOTE improved recall relative to the baseline model by generating synthetic malignant samples. However, SMOTE alone did not sufficiently address misclassification costs, and it sometimes introduced noise, slightly reducing precision.

5.2 Limitations of Cost-Sensitive Learning Alone

Cost-sensitive XGBoost increased sensitivity toward the minority class by penalizing misclassification errors more heavily.

Although this method reduced false negatives, it led to a moderate increase in false positives, reflecting a trade-off between sensitivity and specificity.

5.3 Superiority of the Hybrid Model

The hybrid approach consistently outperformed the individual techniques across all evaluation metrics.

This improvement stems from combining:

- **Enhanced data representation** (via SMOTE), and
- **Misclassification penalization** (via cost-sensitive learning).

The hybrid pipeline produced the **lowest false-negative rate**, the **highest recall**, and the **best AUC-ROC score** among all models.

These metrics indicate that the proposed method effectively strengthens the model's ability to detect malignant cases without excessively compromising precision.

5.4 Clinical Implications

From a medical standpoint, minimizing false negatives is significantly more important than reducing false positives.

A missed malignant diagnosis may lead to delayed treatment and worse patient outcomes.

Therefore, the hybrid model aligns well with clinical decision-support objectives, offering a reliable and ethically sound approach to diagnostic prediction.

6. Conclusion

This research demonstrates that addressing class imbalance in medical datasets requires more than a single technique.

While SMOTE enhances the representation of minority classes and cost-sensitive learning adjusts the decision boundary, neither technique alone is fully sufficient for high-stakes diagnostic tasks.

The proposed **Hybrid SMOTE + Cost-Sensitive XGBoost model** achieved the following:

- The **highest recall**

- The **lowest false-negative count**
- Superior F1-score and AUC-ROC performance
- Better balance between sensitivity and specificity

These results confirm that combining oversampling with algorithmic cost adjustments provides a more robust and clinically meaningful predictive model for breast cancer diagnosis.

Future work may integrate feature selection, ensemble techniques, or advanced deep learning-based methods to further enhance diagnostic reliability.